

Distributed Systems: Google App Engine 2

Stefan Walraven, Wouter Joosen
iMinds-DistriNet, KU Leuven

Overview

- Recap: PaaS
- Google App Engine Services
 - Storage
 - Task Queues
 - Other

Recap: PaaS

- **Platform as a Service (PaaS)**

“A cloud service model that provides the consumer the capability to **deploy** onto the cloud infrastructure **consumer-created or acquired *applications* created using programming languages, libraries, services, and tools supported by the provider.**¹ The **consumer does not manage or control the underlying cloud infrastructure** including network, servers, operating systems, or storage, but has control over the deployed applications and possibly configuration settings for the application-hosting environment.”

[The NIST Definition of Cloud Computing]

¹ This capability does **not necessarily** preclude the use of **compatible** programming languages, libraries, services, and tools from other sources.

Recap: PaaS (2)

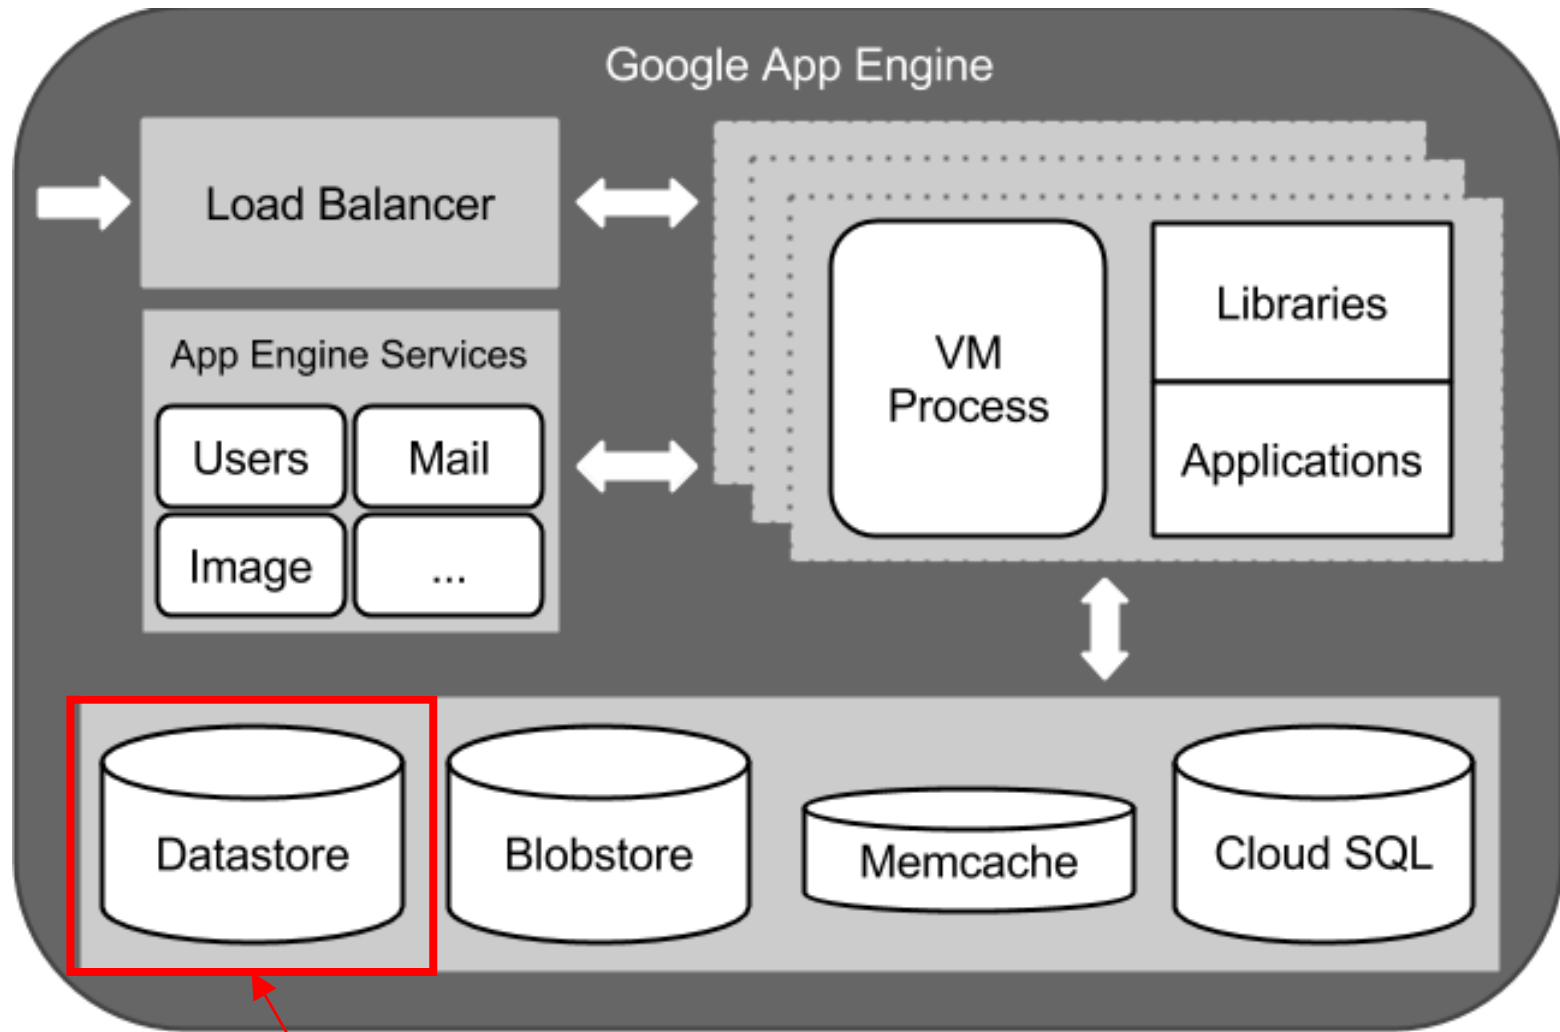
- **Short version:**

PaaS provides a **computing platform and solution stack** upon which applications and services can be **developed and hosted** using programming concepts and tools supported by the provider.

- **Examples:**

- Amazon Elastic Beanstalk
- Cloud Foundry
- Force.com
- Google App Engine
- Oracle PaaS
- Red Hat OpenShift
- Windows Azure

Recap: GAE



Google App Engine Services

■ Storage

- Datastore (subject of GAE session 1)
- Blobstore
- Memcache
- Cloud SQL

■ Task Queues

- Subject of GAE session 2

■ Other:

- Users
- Scheduled tasks
- Fetch URL
- Namespaces
- Mail
- Modules

App Engine Storage

Datastore

- See GAE 1 slides

Blobstore

- **Blobstore API**
 - Cf. Amazon S3 and Windows Azure Blob Storage
 - **Blob** = binary large objects
 - Stored in GAE datastore
 - Goal: serve data objects >> datastore entities
 - created by uploading a file (not via application code)
 - E.g. video or image files

Memcache

■ Memcache API

- **distributed in-memory key-value storage** (not persistent!)
- 2 purposes
 - To speed up common datastore queries (i.e. caching)
 - To store temporary data

■ Low-level API:

```
MemcacheService cache = MemcacheServiceFactory.getMemcacheService();  
value = (<cast to appropriate type>) cache.get(key); // read from cache  
if (value == null) {  
    // get value from other source (e.g. datastore)  
    // ...  
    cache.put(key, value); // populate cache  
}
```

- Supports JCache (JSR 107) (in development)
 - Proposed interface standard for memory caches

Cloud SQL

- Google Cloud SQL
 - Instance == Relational database server
 - Can contain multiple databases
 - Compatible with MySQL
 - Supporting JDBC, JPA, JDO, Hibernate...
 - Configure it correctly! (e.g. connections)
 - Replicated to ensure availability and backup
 - Not free

Other storage services

- Google Cloud Storage (GCS)
 - Storing and service large files
 - Strongly consistent
 - Can also be accessed via Blobstore API
 - Difference with Blobstore unclear
 - GCS seems to be more programmable

App Engine task queues

Message Queues

- Asynchronous inter-process/-thread communication
- Use queue to pass control or content
- ⇔ request-reply
- Examples:
 - Java Message Service (JMS)
 - RabbitMQ
 - JBoss HornetQ

Message Queues (2)

- Often used in cloud computing:
 - to improve performance and scalability
 - by decoupling “bottleneck” activities to worker processes
 - Asynchronous execution
 - No delay at front end (user experience)
 - Opportunity to scale and parallelize execution
 - Increase number of workers
 - Examples:
 - Amazon Simple Queue Service (SQS)
 - Windows Azure Queues
 - **GAE Task Queues**

GAE Task Queues: Queue

- Perform background processing
 - Outside context of user request (asynchronously)
 - Steps:
 - Split up work into *tasks* (== small, discrete unit of work)
 - Insert tasks into *queue* to be executed later
- 2 queue configurations:
 - *Push queue* (default)
 - Automatic scaling of processing capacity to match queue configuration
 - Automatically deletes tasks after processing
 - *Pull queue*
 - More control over when to process tasks
 - Manual scaling and deleting (handled by application)
 - Supports integration with non-App-Engine code

GAE Task Queues: Queue (2)

■ Default (push) queue

- Available for all applications
- No configuration required
- Throughput rate: 5 task invocations per second

```
Queue queue = QueueFactory.getDefaultQueue();
```

■ Custom queues (push & pull)

- Define configuration in **queue.xml** (or **queue.yaml**)
 - name, processing rate, max concurrency, number of retries...
 - in WEB-INF directory inside the WAR
- To group similar types of tasks in same queue

```
Queue queue = QueueFactory.getQueue("<queue name>");
```

GAE Task Queues: Task

- Task
 - Unit of work to be performed by the application
 - Object of `TaskOptions` class
 - Contains an endpoint
 - With request handler for task (i.e. worker)
 - with data payload to parameterize the task (optional)
 - Has a name (optional)
 - Generally unique
 - To delete specific tasks
- 2 kinds of workers:
 - **Default using URL (e.g. to Java Servlet)**
 - Backends (no restrictions, but extra cost per hour per instance)

GAE Task Queues: Task (2)

- Creating and processing tasks:
 - Steps:
 1. Create task and insert in queue
 2. Workers (i.e. task consumer) lease tasks:
 - When task is leased, then unavailable for other workers until lease expires
 3. Workers process leased tasks
 - If success => delete task from queue
 - If lease expires => stop task
 - If failure before lease expires => retry (max. # retries in config)
 - Different depending on queue type:
 - Push queues:
 - automatically by App Engine (lease = 10 minutes)
 - Pull queues:
 - programmatically=> pull queue when worker is available

GAE Task Queues: Push

- Push queues

- Create task:

// call request handler at URL /path_to_worker with parameter key

```
queue.add(withUrl("/path_to_worker").param("key", <value (byte[]/String)>));
```

OR `.payload(<value (byte[]/String)>));`

// OR

// call default request handler at URL /_ah/queue/queue_name

```
queue.add();
```

// OR

// asynchronous call to queue

```
queue.addAsync() / addAsync(<TaskOptions>);
```

- Process task

GAE Task Queues: Push (2)

- Push queues

- Create task
- Process task:
 - Create endpoint for URL (e.g. Java Servlet)

```
public class Worker extends HttpServlet {  
    // default handler method is POST  
    public void doPost(<request>, <response>) {  
        String param = <request>.getParameter("key");  
        OR      InputStream payload = <request>.getInputStream();  
        // unmarshall payload  
        // execute task  
    }  
}
```

- Specify endpoint in **web.xml**
 - Map class name of **Worker** to a url pattern `path_to_worker` (cf. other servlets)

GAE Task Queues: Push (3)

- For each distinct task:
 - Create handler
 - Serialize and deserialize complex arguments
 - Only string or byte[] as payload
 - Cumbersome, especially for many diverse and small tasks
- Solution: **DeferredTask**
 - Interface to define a task as a single method
 - run() method will be called when received (implements **Runnable**)
 - Uses Java serialization (implements **Serializable**)
 - Return == success
 - Exception == failure

GAE Task Queues: Pull

■ Pull Queues

- Create task:

```
queue.add(withMethod(PULL).payload(<payload(byte[]/String)>)  
        .taskName("<name>"));
```

- Lease task (worker inside GAE):

```
List<TaskHandle> tasks =  
    queue.leaseTasks(<# lease units>, <TimeUnit of lease>,  
                    <max # tasks to lease>);
```

- Lease task (worker outside GAE): REST API
- Beta: lease tasks based on tag (i.e. filter tasks)
- Delete task:

```
queue.deleteTask("<name>");
```

GAE Task Queues: Varia

- Securing URLs for Tasks:
 - Prevent users from accessing URLs of workers
 - Restrict access to administrator accounts via `web.xml`
- Push Queues and Development Server
 - Development server
 - Does not respect `<rate>` and `<bucket-size>`
 - Does not retry tasks
 - Does not preserve queue state across server restarts
 - Examine and manipulate tasks from developer console:
 - http://localhost:8888/_ah/admin/taskqueue
 - Works only after one task has been inserted

Other App Engine Services

Users

- Authentication (Users service API) of users with:
 - Google Account
 - Account on Google Apps Domain
 - OpenID identifier (deprecated)

```
UserService userService = UserServiceFactory.getUserService();
```

```
User user = userService.getCurrentUser();
```

Scheduled Tasks

- Scheduled tasks
 - Cron jobs
 - For example, send report email or update cached data
 - 20 scheduled tasks for free
 - At defined times or regular intervals
 - Automatically triggered by GAE Cron Service
 - **cron.xml** OR **cron.yaml** in WEB-INF folder within the WAR package
 - Must define *url* and *schedule*

Fetch URL

- Enables GAE applications to
 - Communicate with other applications
 - Access other resources on the web
- Issue HTTP and HTTPS requests and receive responses
 - Using `java.net.URLConnection` Java standard library
 - Using low-level API
- Actually *the only built-in support for integration* with other applications
 - Other integration support should be implemented by the developer or added via third-party libraries
 - Web services and RESTful web services

Namespaces

■ Namespaces API

■ *Easy partitioning of data*

- Each set of data == namespace
 - Should be set manually (e.g. at the start of a request)
- All namespace-enabled APIs (datastore, memcache and task queues) use the current namespace by default

■ Useful to achieve *multi-tenancy* at the application level

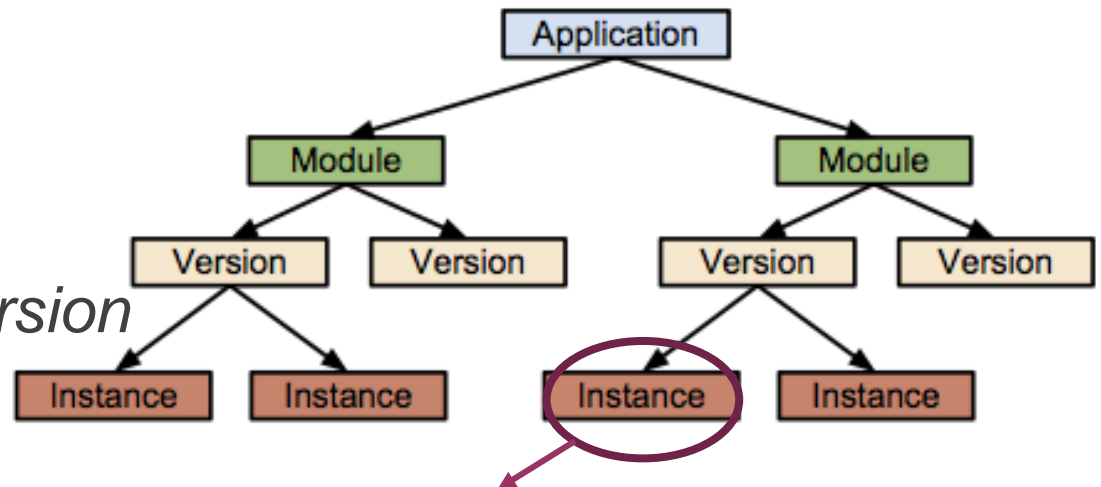
- **Application-level multi-tenancy:** Architectural style that enables a *single application instance* on top of shared hardware and software to serve end users from different *tenants* (i.e. client organizations) simultaneously
 - Key enabler for economies of scale
- Namespace = data partition for one tenant

Mail

- Send email messages on behalf of
 - the application's administrators
 - users with Google Accounts (and are signed in)
 - Retrieve via Users API
- Receive emails at various addresses
- Via JavaMail API (javax.mail)
 - Uses low-level Mail service API
 - No SMTP configuration required/possible

Modules

- Factor large applications into logical components
 - that share stateful services (e.g. Memcache, Datastore, Task Queues)
 - and communicate in secure fashion
 - Cf. *microservices* design pattern
- Each module:
 - Source code
 - Config files
 - Represents a *version*



Separate executable

Modules (2)

- Application using modules is organized as EAR directory structure (cf. Java EE):
 - META-INF directory
 - appengine-application.xml (general info)
 - application.xml (list of modules)
 - Separate WAR subdirectory per module
 - appengine-web.xml (configuration of module)
 - web.xml
- Communication between modules:

```
ModulesService modulesAPI = ModulesServicesFactory.getModulesService();  
String currentModuleName = modulesAPI.getCurrentModule();  
int currentInstance = modulesAPI.getCurrentInstance();  
URL url = new URL("http://" + modulesAPI.getVersionHostname("my-backend-  
module","v1") + "/<web page or application name>");
```


References

- **GAE Developer's Guide**
<https://cloud.google.com/appengine/docs/>
- **Java Service APIs**
<https://cloud.google.com/appengine/docs/java/apis>
- **Task Queue Java API Overview** ←
<https://cloud.google.com/appengine/docs/java/taskqueue/>
- **Using Push Queues** ←
<https://cloud.google.com/appengine/docs/java/taskqueue/overview-push>
- **Java Task Queue Configuration**
<https://cloud.google.com/appengine/docs/java/config/queue>
- **Java Task Queue Javadoc**
<https://cloud.google.com/appengine/docs/java/javadoc/com/google/appengine/api/taskqueue/package-summary>